

BITS PILANI WILP

Software Engineering for Machine Learning (AIMLCZG546)

Assignment II — Implementation, Code Quality & Quality Assurance

Project: Loan Approval Risk Service — automated consumer credit underwriting
Group Number: 84 · **Date:** August 2026 · **Builds on:** Assignment I (Group 84)

Group Members & Contributions

Sl.	BITS ID	Name	Qualitative Contribution	%
1	2025AA05710	Singh Pritesh	Refactoring to the <code>loan_risk</code> package, error handling & structured logging, lint/format toolchain, quality gates	100
2	2025AA05368	Gangera Tushar	Research-vs-production analysis, data-quality metrics, schema contract and data-validation test suite	100
3	2025AB05154	Gangam Shuba Nandini	ML behavioural tests (training & inference), model-quality metrics, calibration and drift monitoring	100
4	2025AA05574	Shaifali Garg	FastAPI design & implementation, integration tests, production experimentation and security analysis	100

Executive Summary

Assignment I delivered a working prototype: a four-filter pipe-and-filter pipeline behind a single FastAPI endpoint. This assignment turns that prototype into a system that could be handed to an operations team. The scoring logic was rebuilt as an installable package (`src/loan_risk/`) with one cohesive class per responsibility; a research notebook is retained unchanged beside its engineered counterpart so the difference is evidence rather than assertion; error handling and structured JSON logging were added across the ingestion, validation and training paths; the whole codebase was put under `isort`, `black`, `flake8` and `pylint`; and the API was redesigned with versioned resources, a bounded batch endpoint and differentiated status codes.

On the QA side, 84 pytest tests were written across four distinct test types, including the ML-specific families the assignment calls for — overfit-a-small-batch, loss-decreases, output shape/range, directional and invariance. Four model-quality metrics and four data-quality metrics are measured and gated. The measured headline numbers are: accuracy **0.9480**, F1 **0.9343**, ROC-AUC **0.9881**, Brier **0.0458**, mean latency **10.29 ms** against a 150 ms SLA, flake8 **45** → **0** violations and pylint **10.00/10**.

The test suite earned its keep before submission. The invariance test `test_prediction_is_invariant_to_repeated_calls` failed on the first run: with `n_jobs=-1` the forest accumulated tree votes in non-deterministic thread order, so two identical requests could differ by $\sim 5e-17$ — enough to flip an applicant sitting exactly on the 0.50 cut-off between APPROVED and DENIED. Section 7.3 documents the defect and the fix.

Objective 1 — Implementation and Code Sharing

1. Refactoring with OOP and Functional Principles

The Assignment I code was a flat app/ folder in which pipeline.py held four module-level functions that each reached into a global settings object, and main.py called joblib.load directly. That is serviceable for a demo and awkward to test, extend or reason about. The refactor splits the system along responsibility boundaries, one module per concern:

Module	Type	Responsibility	Paradigm
config.py	Settings et al.	Frozen dataclasses loaded from YAML; every threshold is data	OOP (immutable value objects)
data/ingestion.py	DataIngestor	The only component that reads training data from disk	OOP
data/validation.py	DataValidator	Declarative schema contract; emits DataQualityReport	OOP + declarative spec
features/engineering.py	FeatureEngineer	Derived ratios; sklearn transformer wrapping pure functions	Functional core, OOP shell
models/trainer.py	ModelTrainer	train → evaluate → gate → persist lifecycle	OOP
models/predictor.py	ModelRegistry, RiskPredictor	Artifact custody; the four scoring filters	OOP + pure filters
monitoring/drift.py	DriftMonitor	PSI and KS drift against the frozen reference profile	Functional core, OOP shell
api/	schemas.py, app.py	HTTP contract, routing and exception→status mapping	Declarative (Pydantic)

Figure 1: Module responsibility map of the loan_risk package.

Where each paradigm was chosen, and why

- **Object-oriented** where there is genuine state or a lifecycle to own: ModelRegistry owns the loaded artifact, ModelTrainer owns the fitted pipeline and its metrics. Constructor injection of Settings replaced the global lookups, so any test can substitute a different configuration without patching a module.
- **Functional** where the operation is a transformation: safe_ratio, compute_loan_to_income, assign_risk_tier and validate_business_rules are pure — same input, same output, no mutation. test_validation_filter_does_not_mutate_the_payload and test_transformer_does_not_mutate_its_input assert that purity rather than trusting it.
- **Immutability** for configuration: the Settings aggregate is a frozen dataclass, so a request handler cannot quietly rewrite the decision threshold for every other request in the process.
- **Composition over duplication**: FeatureEngineer is a scikit-learn TransformerMixin and is therefore the *first step of the serialised pipeline*. Training and serving cannot disagree about feature definitions, because they execute the same pickled object.

The last point is the single most valuable structural change. In Assignment I the derived ratios were computed in data/prepare_data.py for training and again in extract_features() for serving — two copies of one definition, which is the textbook setup for training/serving skew. There is now one copy, and it travels inside the artifact.

```
if algorithm == "random_forest":
    estimator = RandomForestClassifier(
        n_estimators=model_cfg.n_estimators,
        max_depth=model_cfg.max_depth,
        min_samples_leaf=model_cfg.min_samples_leaf,
        random_state=model_cfg.random_state,
        # n_jobs=1 is a *correctness* choice, not a performance
        # oversight. With n_jobs=-1 the per-tree vote accumulation
        # happens in non-deterministic thread order, so two identical
        # requests can differ at ~1e-16 -- enough to flip an applicant
        # sitting exactly on the 0.50 cut-off. Regulated lending
        # decisions must be bit-reproducible for audit, and inference
        # is already ~2 ms per request, so there is nothing to buy.
```

```

        # Caught by tests/test_model_inference.py::
        # test_prediction_is_invariant_to_repeated_calls.
        n_jobs=1,
    )
    steps = [("features", FeatureEngineer()), ("model", estimator)]
    ... "pipeline_built", extra={"algorithm": algorithm, "n_steps": len(steps)}
    )
    return Pipeline(steps)

```

Listing 1: the estimator is assembled as a Pipeline whose first stage is the serving-time feature transformer.

2. Research Code vs Production Code

The component chosen for this comparison is **feature engineering**. Both artefacts are in the repository and both compute the same two ratios:

- **Research:** notebooks/research_prototype.ipynb (also exported verbatim as legacy/research_feature_prototype.py so the linters can be pointed at it).
- **Production:** src/loan_risk/features/engineering.py.

The research notebook is not a strawman — it is what genuinely useful exploratory code looks like. It answered the question it was written for. The point of the table below is that every one of its shortcuts is a defect only once the code has to run unattended, and each has a specific engineered answer.

#	Research code (notebook)	Consequence in production	Production answer
1	Hard-coded absolute path C:\Users\analyst\Desktop /...	Runs on exactly one machine	Settings.path() resolves every path from config.yaml
2	train_test_split with no random_state	Metrics move every run; results are unauditable	random_state from config; test_training_is_reproducible asserts bit-identical output
3	LoanAmount/AnnualIncome with no guard	inf / ZeroDivisionError on a zero-income application	safe_ratio() applies +1 smoothing; a finiteness check raises FeatureEngineeringError
4	Feature list retyped by hand in the notebook	Training/serving skew as soon as one copy changes	One FEATURE_ORDER contract, sourced from config and re-imposed on every transform
5	Dead variables (tmp, ratio2, x)	Reader cannot tell what is load-bearing	Every symbol is used; flake8 F401/F841 enforce it
6	except: pass	Failures are silent	Typed exception hierarchy; every handler logs at a level and re-raises
7	No logging, no types, no docstrings	Nothing is observable; behaviour is undocumented	Structured JSON logging, full type hints, docstring on every public callable (pylint 10.00/10)
8	Logic lives in notebook cells	Cannot be imported, reused or unit-tested	Importable package; 14 unit tests cover this module alone
9	print() for diagnostics	Unstructured output, unusable by a log aggregator	logger.info(..., extra={...}) emitting one JSON object per event

Figure 2: Defect-by-defect mapping from the research notebook to its production counterpart.

Side-by-side: the same computation, twice

Research — notebooks/research_prototype.ipynb

```

df['LoanToIncomeRatio'] = df['LoanAmount']/df['AnnualIncome'] # NOTE: blows up when income is 0
df['SavingsToLoanRatio'] = df['SavingsAccountBalance']/df['LoanAmount']
df['ratio2'] = df['MonthlyDebtPayments']*12/df['AnnualIncome']
df['x'] = df['CreditScore']/850
FEATURES=['Age', 'AnnualIncome', 'CreditScore', 'LoanAmount', 'LoanToIncomeRatio', ...]

def get_ratio(a,b) :
    return a/b # no zero guard, no types, no docstring

```

Production — src/loan_risk/features/engineering.py

```

def safe_ratio(
    numerator: pd.Series | float, denominator: pd.Series | float
) -> pd.Series | float:
    """Divide with +1 smoothing so the denominator can never be zero.

    Pure function -- no I/O, no globals -- which is exactly why it can be
    property-tested in isolation (see ``tests/test_unit_features.py``).
    """
    return numerator / (denominator + _EPSILON)

def compute_loan_to_income(loan_amount, annual_income):
    """Leverage ratio: how many times the annual income is being borrowed."""
    return safe_ratio(loan_amount, annual_income)

def compute_savings_to_loan(savings_balance, loan_amount):
    """Cushion ratio: liquid savings held against the requested principal."""
    return safe_ratio(savings_balance, loan_amount)

def transform(self, X: pd.DataFrame) -> pd.DataFrame: # noqa: N803
    """Return a frame containing exactly ``feature_order``, in order.

    Raises:
        FeatureEngineeringError: if a required base column is absent or the
        derived values are not finite.
    """
    if not isinstance(X, pd.DataFrame):
        raise FeatureEngineeringError(
            f"Expected a pandas DataFrame, received {type(X).__name__}"
        )

    frame = X.copy()
    missing = [c for c in BASE_COLUMNS if c not in frame.columns]
    if missing:
        logger.error("feature_base_columns_missing", extra={"missing": missing})
        raise FeatureEngineeringError(f"Missing base columns: {missing}")

    frame["LoanToIncomeRatio"] = compute_loan_to_income(
        frame["LoanAmount"], frame["AnnualIncome"]
    )
    frame["SavingsToLoanRatio"] = compute_savings_to_loan(
        frame["SavingsAccountBalance"], frame["LoanAmount"]
    )

    derived = frame[["LoanToIncomeRatio", "SavingsToLoanRatio"]]
    if not np.isfinite(derived.to_numpy(dtype=float)).all():
        logger.error("derived_features_not_finite")
        raise FeatureEngineeringError(
            "Derived ratios contain NaN/inf -- check inputs for nulls."
        )

    out_of_range = int((frame["LoanToIncomeRatio"] > 5.0).sum())
    if out_of_range:
        # Recoverable: the model still scores these, but the operator should
        # know that unusually leveraged applications are arriving.
        logger.warning(
            "high_leverage_applications_detected",
            extra={"count": out_of_range, "threshold_loan_to_income": 5.0},
        )

    logger.info(
        "feature_engineering_completed",
        extra={"rows": int(len(frame)), "n_features": len(self.feature_order)},
    )
    return frame[self.feature_order]

```

Listing 2: the production transformer — declared contract, defensive division, explicit failure modes, structured logging, and a sklearn interface so it can be serialised with the model.

3. Error Handling and Logging

A typed exception hierarchy rooted at `LoanRiskError` lets the API layer tell our failures apart from unexpected ones and map each to the right status code, instead of returning a blanket 500. Logging is structured JSON on stdout — the same records that help a developer locally become the monitoring substrate in production without a second instrumentation path.

Level	Policy in this codebase	Example event
INFO	Normal lifecycle events worth auditing	<code>data_ingestion_completed</code> , <code>model_artifact_saved</code> , <code>loan_application_scored</code>

Level	Policy in this codebase	Example event
WARNING	Recoverable anomaly: a fallback was taken, a soft data rule was breached, or a request was rejected by policy	duplicate_rows_detected, data_quality_violations_detected, business_rule_rejected, data_drift_detected
ERROR	The operation failed and the caller receives an error	data_source_missing, feature_contract_violation, quality_gate_breached, model_training_failed

Figure 3: Log-level policy, applied consistently across all modules.

The three critical functions

Critical function 1 — `DataIngestor.load()`: every I/O failure mode (missing file, empty file, unparseable CSV, zero rows) is caught, logged and re-raised as a single `DataIngestionError`, so callers never have to catch `OSError`. Duplicate rows are a **WARNING** — the run continues, but the operator is told.

```
def load(self, path: Path | str | None = None) -> pd.DataFrame:
    """Read a CSV into a dataframe.

    Raises:
        DataIngestionError: file missing, empty, or unparseable.
    """
    source = Path(path) if path else self.config.path("processed_data")
    logger.info("data_ingestion_started", extra={"source": str(source)})

    if not source.exists():
        logger.error(
            "data_source_missing",
            extra={"source": str(source), "hint": "run scripts/build_dataset.py"},
        )
        raise DataIngestionError(f"Dataset not found at {source}")

    try:
        frame = pd.read_csv(source)
    except pd.errors.EmptyDataError as exc:
        logger.error("data_source_empty", extra={"source": str(source)})
        raise DataIngestionError(f"Dataset at {source} is empty") from exc
    except (pd.errors.ParserError, UnicodeDecodeError, OSError) as exc:
        logger.error(
            "data_source_unreadable",
            extra={"source": str(source), "error": str(exc)},
        )
        raise DataIngestionError(f"Could not parse {source}: {exc}") from exc

    if frame.empty:
        logger.error("data_source_zero_rows", extra={"source": str(source)})
        raise DataIngestionError(f"Dataset at {source} contains zero rows")

    duplicates = int(frame.duplicated().sum())
    if duplicates:
        # Recoverable: we keep going but the operator must know.
        logger.warning(
            "duplicate_rows_detected",
            extra={"duplicate_rows": duplicates, "total_rows": len(frame)},
        )

    logger.info(
        "data_ingestion_completed",
        extra={"rows": int(frame.shape[0]), "columns": int(frame.shape[1])},
    )
    return frame
```

Critical function 2 — `DataValidator.validate()`: soft breaches accumulate and log at **WARNING**; in `strict=True` mode (used by the training entrypoint) a breach logs at **ERROR** and raises `SchemaValidationError` carrying the full violation list. Training on silently corrupt data is the failure this prevents.

```
def validate(self, frame: pd.DataFrame, strict: bool = False) -> DataQualityReport:
    if frame is None or frame.empty:
        logger.error("validation_called_on_empty_frame")
        raise SchemaValidationError("Cannot validate an empty dataframe.")

    violations: List[str] = []
    conforming_columns = 0
    for spec in self.schema:
        findings = self._check_column(frame, spec)
        violations.extend(findings)
        conforming_columns += int(not findings)
```

```

# ---- DQ metric 1: schema conformance rate -----
conformance = conforming_columns / len(self.schema)

# ---- DQ metric 2: missing-value fractions -----
null_fractions = frame.isna().mean()
overall_missing = float(frame.isna().to_numpy().mean())
worst_column = str(null_fractions.idxmax()) if len(null_fractions) else None
worst_fraction = float(null_fractions.max()) if len(null_fractions) else 0.0

if overall_missing > self.max_missing_fraction:
    violations.append(
        f"missing_fraction_exceeded:{overall_missing:.4f}"
        f">{self.max_missing_fraction}"
    )

report = DataQualityReport(
    n_rows=int(frame.shape[0]),
    n_columns=int(frame.shape[1]),
    schema_conformance_rate=conformance,
    missing_value_fraction=overall_missing,
    worst_column_missing_fraction=worst_fraction,
    worst_column=worst_column,
    violations=violations,
)

if violations:
    logger.warning(
        "data_quality_violations_detected",
        extra={
            "n_violations": len(violations),
            "violations": violations[:10],
            "schema_conformance_rate": round(conformance, 4),
        },
    )
    if strict:
        logger.error(
            "data_quality_gate_failed",
            extra={"n_violations": len(violations)},
        )
        raise SchemaValidationError(
            f"{len(violations)} schema violation(s) detected.", violations
        )
else:
    logger.info(
        "data_quality_gate_passed",
        extra={
            "rows": report.n_rows,
            "schema_conformance_rate": 1.0,
            "missing_value_fraction": round(overall_missing, 6),
        },
    )
return report

```

Critical function 3 — `ModelTrainer.train()` / `enforce_quality_gates()`: length mismatch and a single-class target are rejected before fitting; a fit exception is wrapped as `ModelTrainingError`; and a model that misses a release gate is logged at `ERROR` and refuses to ship.

```

def enforce_quality_gates(self, metrics: Optional[ModelMetrics] = None) -> None:
    metrics = metrics or self.metrics
    if metrics is None:
        raise ModelTrainingError("No metrics available to gate on.")

    gates = self.config.gates
    breaches = []
    if metrics.accuracy < gates.min_accuracy:
        breaches.append(f"accuracy {metrics.accuracy:.4f} < {gates.min_accuracy}")
    if metrics.f1 < gates.min_f1:
        breaches.append(f"f1 {metrics.f1:.4f} < {gates.min_f1}")
    if metrics.roc_auc < gates.min_roc_auc:
        breaches.append(f"roc_auc {metrics.roc_auc:.4f} < {gates.min_roc_auc}")
    if metrics.brier_score > gates.max_brier_score:
        breaches.append(f"brier {metrics.brier_score:.4f} > {gates.max_brier_score}")

    if breaches:
        logger.error("quality_gate_breached", extra={"breaches": breaches})
        raise ModelTrainingError("Quality gates breached: " + "; ".join(breaches))
    logger.info("quality_gates_passed", extra=metrics.as_dict())

```

Sample of the real emitted log stream (reports/metrics/training_log.txt):

```

{"timestamp": "2026-08-03T22:33:34", "level": "INFO", "logger": "loan_risk.data.ingestion", "message": "data_ingestion_c"}
{"timestamp": "2026-08-03T22:33:34", "level": "INFO", "logger": "loan_risk.data.validation", "message": "data_quality_ga"}
{"timestamp": "2026-08-03T22:33:40", "level": "INFO", "logger": "loan_risk.models.trainer", "message": "model_evaluated"}
{"timestamp": "2026-08-03T22:33:41", "level": "WARNING", "logger": "loan_risk.monitoring.drift", "message": "data_drift"}
{"timestamp": "2026-08-03T22:33:41", "level": "INFO", "logger": "loan_risk.models.predictor", "message": "loan_applicati"}

```

Listing 3: every record is one JSON object, so `latency_ms` or `risk_tier` can be aggregated directly by ELK/Splunk with no log parsing.

4. Code Formatting and Linting

Four tools are configured, each with a distinct job: `isort` orders imports, `black` owns formatting outright, `flake8` catches correctness and complexity issues formatting cannot, and `pylint` adds design-level checks. Configuration lives in `pyproject.toml` (`black`, `isort`, `pytest`, `pylint`) and `setup.cfg` (`flake8`, which still has no `pyproject` support). The line length is set to 90 in all four so they cannot fight each other.

Method

The **before** snapshot was taken on the codebase as first written, together with the untouched research code in `legacy/`. It reported **45 flake8 violations** across the tree and **13 files** that `black` would reformat. The formatters were then run, and the two findings the formatters cannot fix were repaired by hand — one of which was a genuine design issue (C901: `DataValidator.validate` had a cyclomatic complexity of 11 against a budget of 10, and was split into `_check_column` and `_check_range`). Both reports are stored verbatim in `reports/lint/`.

Check	Command	Before	After
flake8 violations	<code>python -m flake8 src scripts tests legacy</code>	45	0
black — files needing reformat	<code>python -m black --check src scripts tests</code>	13	0
isort — import order	<code>python -m isort --check-only src scripts tests</code>	clean	clean
pylint score	<code>python -m pylint src/loan_risk</code>	9.72 / 10	10.00 / 10

Figure 4: Lint and format results, before and after. Raw reports: `reports/lint/01_before_flake8.txt` → `07_after_flake8.txt`.




Figure 5: The same data as a chart. The pylint score is scaled ×10 to share the axis.

Representative violations from the BEFORE report

```

legacy\research_feature_prototype.py:5:1: F401 'numpy as np' imported but unused
legacy\research_feature_prototype.py:5:20: E401 multiple imports on one line
legacy\research_feature_prototype.py:6:1: F401 'os' imported but unused
legacy\research_feature_prototype.py:6:1: F401 'sys' imported but unused
legacy\research_feature_prototype.py:6:1: F401 'json' imported but unused
legacy\research_feature_prototype.py:6:10: E401 multiple imports on one line
legacy\research_feature_prototype.py:12:91: E501 line too long (93 > 90 characters)
legacy\research_feature_prototype.py:14:7: E201 whitespace after '('
legacy\research_feature_prototype.py:14:17: E202 whitespace before ')'
legacy\research_feature_prototype.py:18:27: E231 missing whitespace after ','
legacy\research_feature_prototype.py:22:91: E501 line too long (101 > 90 characters)
legacy\research_feature_prototype.py:26:9: E225 missing whitespace around operator
legacy\research_feature_prototype.py:26:16: E231 missing whitespace after ','
legacy\research_feature_prototype.py:26:31: E231 missing whitespace after ','

```


Listing 4: extract from reports/lint/01_before_flake8.txt. E401 multiple imports per line, F401 unused import, E722 bare except, E231 missing whitespace, C901 too complex.

AFTER — the same command on the production tree

```
$ python -m flake8 src scripts tests
$ python -m black --check src scripts tests
All done! 25 files would be left unchanged.
$ python -m isort --check-only src scripts tests
$ python -m pylint src/loan_risk
```

Your code has been rated at 10.00/10 (previous run: 9.72/10, +0.28)

Listing 5: flake8 and isort print nothing on success. legacy/ is deliberately excluded from black and isort — it is the evidence, and reformatting it would destroy the comparison.

5. REST API Design and Implementation

The service is a FastAPI application. Assignment I exposed a single unversioned /predict; the redesign applies the API-design practices the assignment asks for.

- **Versioned, resource-oriented paths.** /v1/predict, /v1/predict/batch, /v1/model/metadata. A breaking change can ship as /v2 without stranding callers. /health stays unversioned because orchestrator probes should never be tied to an API generation.
- **Explicit request/response schemas.** Every field is bounded with ge/le, and extra="forbid" rejects unknown keys — so a client typo like credit_scr is a loud 422 rather than a silently defaulted feature.
- **Correct, differentiated status codes.** 200 scored · 422 schema violation · 400 business-rule rejection · 503 model unavailable · 404 unknown route. A policy rejection is a client error and must not be reported as a server fault.
- **A declared response_model on every route,** so the OpenAPI contract is generated from the code and cannot drift from it.
- **A bounded batch endpoint** (1–100 applications). The cap is a denial-of-service control as much as an ergonomic one.
- **Exception handlers, not try/except in routes.** Domain exceptions are translated centrally, so internal messages and stack traces never reach the client — asserted by test_error_responses_never_leak_internals.
- **Graceful degradation.** If the artifact is missing the process still starts, /health reports degraded, and scoring returns 503 — so an orchestrator keeps the previous replica serving instead of crash-looping.

REST API contract generated from the code (GET /openapi.json)

Method	Path	Summary	Tag	Documented status codes
GET	/health	Health	ops	200
GET	/v1/model/metadata	Model Metadata	ops	200
POST	/v1/predict	Predict	inference	200, 400, 422, 503
POST	/v1/predict/batch	Predict Batch	inference	200, 400, 422, 503

Figure 6: The endpoint table, generated from the live service's own /openapi.json.

```
@app.get("/health", response_model=HealthResponse, tags=["ops"])
def health() -> HealthResponse:
    """Liveness/readiness probe for Kubernetes."""
    return HealthResponse(
        status="healthy" if registry.is_loaded else "degraded",
        model_loaded=registry.is_loaded,
        model_version=registry.version,
        environment=settings.env,
        api_version=settings.version,
    )

def predict(application: LoanApplication) -> RiskAssessmentResponse:
    """Score a single loan application."""
    assessment = predictor.predict(application.to_model_payload())
```



```
return RiskAssessmentResponse(**assessment.as_dict())
```

Listing 6: route definitions — documented status codes, declared response models, tags for the Swagger grouping.

Verification against a running service

The figures below are the recorded transcript of real HTTP calls issued against a live uvicorn process on 127.0.0.1:8077 (reports/metrics/api_transcript.json).

Live service: health probe and two scored applications (HTTP 200)

GET /health <i>Health probe</i> request (no request body) response <pre>{ "status": "healthy", "model_loaded": true, "model_version": "2.0.0", "environment": "production", "api_version": "2.0.0" }</pre>	HTTP 200
POST /v1/predict <i>Low-risk applicant (200)</i> request <pre>{ "age": 45, "annual_income": 95000, "credit_score": 720, "employment_status": 0, "education_level": 4, "loan_amount": 25000, "loan_duration": 36, "monthly_debt_payments": 400, "credit_card_utilization_rate": 0.25, "debt_to_income_ratio": 0.15, "bankruptcy_history": 0, "loan_purpose": 3, "previous_loan_defaults": 0, }</pre> response <pre>{ "is_approved": true, "probability": 0.939, "risk_tier": "LOW", "net_worth": 220000, "debt_to_income": 0.15, "latency_ms": 8.29, "model_version": "2.0.0" }</pre>	HTTP 200
POST /v1/predict <i>High-risk applicant (200)</i> request <pre>{ "age": 45, "annual_income": 95000, "credit_score": 540, "employment_status": 0, "education_level": 4, "loan_amount": 90000, "loan_duration": 36, "monthly_debt_payments": 400, "credit_card_utilization_rate": 0.92, "debt_to_income_ratio": 0.85, "bankruptcy_history": 0, "loan_purpose": 3, "previous_loan_defaults": 1, }</pre> response <pre>{ "is_approved": false, "probability": 0.0805, "risk_tier": "HIGH", "net_worth": 220000, "debt_to_income": 0.85, "latency_ms": 6.05, "model_version": "2.0.0" }</pre>	HTTP 200

Figure 7: Health probe and two scored applications. The same schema-valid payload with a 720 credit score and a 540 score with prior defaults yields LOW ($p=0.939$, approved) and HIGH ($p=0.081$, denied) respectively.

Live service: differentiated error handling (422 schema, 422 typo, 400 business rule)**POST /v1/predict****HTTP 422***Schema violation: credit_score=1500 (422)*

```
request
{"age": 45, "annual_income": 95000, "credit_score": 1500, "employment_status": 0, "education_level": 4, "loan_amount":
25000, "loan_duration": 36, "monthly_debt_payments": 400, "credit_card_utilization_rate": 0.25,
"debt_to_income_ratio": 0.15, "bankruptcy_history": 0, "loan_purpose": 3, "previous_loan_defaults": 0,
response
{
  "detail": [
    {
      "type": "less_than_equal",
      "loc": [
        "body",
        "credit_score"
      ],
      "msg": "Input should be less than or equal to 850",
      "input": 1500,
      "ctx": {
        ... (truncated)
      }
    }
  ]
}
```

POST /v1/predict**HTTP 422***Unknown field typo (422)*

```
request
{"age": 45, "annual_income": 95000, "credit_score": 720, "employment_status": 0, "education_level": 4, "loan_amount":
25000, "loan_duration": 36, "monthly_debt_payments": 400, "credit_card_utilization_rate": 0.25,
"debt_to_income_ratio": 0.15, "bankruptcy_history": 0, "loan_purpose": 3, "previous_loan_defaults": 0,
response
{
  "detail": [
    {
      "type": "extra_forbidden",
      "loc": [
        "body",
        "credit_scr"
      ],
      "msg": "Extra inputs are not permitted",
      "input": 700
    }
  ]
  ... (truncated)
}
```

POST /v1/predict**HTTP 400***Business-rule breach (400)*

```
request
{"age": 45, "annual_income": 20000, "credit_score": 720, "employment_status": 0, "education_level": 4, "loan_amount":
500000, "loan_duration": 36, "monthly_debt_payments": 400, "credit_card_utilization_rate": 0.25,
"debt_to_income_ratio": 0.15, "bankruptcy_history": 0, "loan_purpose": 3, "previous_loan_defaults": 0,
response
{
  "detail": "Requested amount exceeds 5x annual income.",
  "error_type": "BusinessRuleViolation"
}
```

Figure 8: Differentiated error handling. Out-of-range value → 422 with the offending field named; unknown key → 422 extra_forbidden; schema-valid but over-leveraged → 400 with a business reason and no internal detail.

Objective 2 — Quality Assurance

6. Test Types Implemented

The suite contains **84 tests**, all passing, organised into four distinct types and tagged with pytest markers so each layer can be run independently in CI (pytest -m unit for the fast pre-commit gate, the full suite before merge).

Type	Marker	File	Count	What it proves
Unit	unit	test_unit_features.py	14	Individual pure functions and the transformer behave in isolation: correct values, determinism, no mutation, correct failure modes
Integration	integration	test_integration_api.py	21	Routing → Pydantic → registry → predictor → transformer → estimator → serialisation are wired together correctly, exercised over real HTTP
Data validation	data	test_data_validation.py	15	The data contract holds: schema conformance, missing values, PSI/KS drift, ingestion feature contract
ML behavioural	ml	test_model_training.py, test_model_inference.py	34	Learning actually happens, and inference obeys the domain's shape, range, directional and invariance expectations

Figure 9: Test inventory. Fixtures build every artefact in memory, so the suite runs on a clean checkout with no trained model on disk.

```
$ python -m pytest tests -q
===== 84 passed, 1 warning in 8.81s =====

$ python -m pytest tests -m unit --co -q
14/84 tests collected (70 deselected)
```

Listing 7: full-suite result. Raw output is stored in reports/metrics/pytest_output.txt.

7. Tests for ML Components

7.1 Testing model training

A model that trains without raising is not a model that has learned. Five training-specific tests establish that it has:

Test	Assertion	Defect it catches
test_model_can_overfit_a_small_batch	A high-capacity forest reaches ≥ 0.95 training accuracy on 40 rows	Features misaligned with labels, or the pipeline dropping the signal before the estimator
test_training_loss_decreases_with_capacity	Log-loss falls monotonically over $\text{max_depth} \in \{1, 3, 8, \text{None}\}$, and the last is under half the first	The model is not fitting the batch at all
test_training_is_reproducible	Two runs with the same seed give identical probabilities	Unseeded randomness; unauditable results
test_shuffled_labels_destroys_generalisation	With permuted labels, held-out AUC collapses to 0.35–0.65	Target leakage or a broken evaluation split — if a model can still 'predict' random labels, the metrics are meaningless
test_quality_gates_reject_a_weak_model	A depth-1, one-tree stump trips the release gate	A gate that silently passes everything

Figure 10: Model-training tests.

```
def test_model_can_overfit_a_small_batch(xy):
    """A high-capacity model must memorise 40 rows (accuracy >= 0.95).

    Failure here means labels are misaligned with features, or the pipeline is
    dropping the signal before it reaches the estimator.
    """
    features, labels = xy
    small_x = features.head(40)
```

```

small_y = labels.head(40)

trainer = ModelTrainer(settings)
trainer.train(small_x, small_y, evaluate=False)
train_accuracy = trainer.pipeline.score(small_x, small_y)
assert train_accuracy >= 0.95, f"Could not overfit 40 rows: {train_accuracy:.3f}"

```

Listing 8: the canonical overfit-a-small-batch test.

Measured loss curve produced by the capacity sweep: **0.4290** → **0.2936** → **0.1146** → **0.0940** for max_depth 1 → 3 → 8 → unbounded. Depth, not tree count, is the right capacity axis here: an unconstrained forest already interpolates the batch at a single tree, so sweeping n_estimators would measure variance reduction rather than learning. Our first version of this test made exactly that mistake and failed; the test was wrong, not the model.

7.2 Testing model inference

Inference tests follow the three standard families:

Family	Representative test	Expectation
Shape & range	test_batch_prediction_shape_is_correct test_probabilities_are_valid_and_sum_to_one test_output_stays_in_range_across_the_credit_spectrum	37 rows in → (37, 2) out; every probability in [0, 1] and rows summing to 1; finite output for credit scores 300–850
Directional	test_higher_credit_score_does_not_reduce_approval_probability test_higher_debt_to_income_does_not_increase_approval_probability test_a_much_larger_loan_does_not_increase_approval_probability	Domain monotonicity holds: 520→800 credit score cannot lower P(approve); 0.10→0.95 DTI cannot raise it; a 15k→140k loan cannot raise it
Invariance	test_prediction_is_invariant_to_dictionary_key_order test_prediction_is_invariant_to_column_order test_prediction_is_invariant_to_batching test_prediction_is_invariant_to_repeated_calls	Payload key order, dataframe column order, batched vs row-by-row scoring, and repetition must all leave the score unchanged

Figure 11: Model-inference tests. The column-order test is what proves the FeatureEngineer really does re-impose the feature contract inside the artifact.

```

def test_higher_credit_score_does_not_reduce_approval_probability(predictor):
    poor = _score(predictor, CreditScore=520)
    excellent = _score(predictor, CreditScore=800)
    assert excellent >= poor, (poor, excellent)

def test_higher_debt_to_income_does_not_increase_approval_probability(predictor):
    low = _score(predictor, DebtToIncomeRatio=0.10)
    high = _score(predictor, DebtToIncomeRatio=0.95)
    assert high <= low, (low, high)

def test_a_much_larger_loan_does_not_increase_approval_probability(predictor):
    modest = _score(predictor, LoanAmount=15_000)
    aggressive = _score(predictor, LoanAmount=140_000)
    assert aggressive <= modest, (modest, aggressive)

```

Listing 9: directional tests. They assert the weak inequality — a tree ensemble is not globally monotonic, and demanding strict monotonicity would produce a flaky test that teams learn to ignore.

7.3 A defect the suite actually caught

test_prediction_is_invariant_to_repeated_calls failed on its first run. Scoring the identical payload five times produced four distinct floats. The cause was n_jobs=-1 on the RandomForestClassifier: per-tree vote accumulation happens in whatever order the worker threads finish, and floating-point addition is not associative. The spread was $\sim 5 \times 10^{-14}$ — invisible in any aggregate metric, and completely unacceptable here, because an applicant whose probability sits exactly on the 0.50 cut-off could receive APPROVED on one request and DENIED on an identical retry. In regulated lending, decisions must be bit-reproducible for audit.

The fix was to pin the estimator to a deterministic reduction. Inference is ~10 ms per request against a 150 ms SLA, so there was no performance argument for keeping the parallel path.

```
estimator = RandomForestClassifier(
    n_estimators=model_cfg.n_estimators,
    max_depth=model_cfg.max_depth,
    min_samples_leaf=model_cfg.min_samples_leaf,
    random_state=model_cfg.random_state,
    # n_jobs=1 is a *correctness* choice, not a performance
    # oversight. With n_jobs=-1 the per-tree vote accumulation
    # happens in non-deterministic thread order, so two identical
    # requests can differ at ~1e-16 -- enough to flip an applicant
    # sitting exactly on the 0.50 cut-off. Regulated lending
    # decisions must be bit-reproducible for audit, and inference
    # is already ~2 ms per request, so there is nothing to buy.
    # Caught by tests/test_model_inference.py::
    # test_prediction_is_invariant_to_repeated_calls.
    n_jobs=1,
```

Listing 10: the fix, with the reason recorded at the point of the decision and a pointer back to the test that found it.

8. Model-Quality and Data-Quality Metrics

8.1 Model quality — four metrics

Accuracy alone is a poor summary for underwriting: the class balance is uneven (40.3% approvals) and the *probability*, not just the label, drives the risk tier. Four complementary metrics are therefore measured on a stratified held-out split of 4,000 applications, and all four are enforced as release gates in `config.yaml`.

ID	Metric	Why it is measured	Gate	Measured	Status
MQ-1	Accuracy	Headline decision correctness	≥ 0.80	0.9480	PASS
MQ-2	F1 score	Balances the cost of a missed good customer against an approved bad one under class imbalance	≥ 0.80	0.9343	PASS
MQ-3	ROC-AUC	Threshold-independent ranking quality; survives a change to the 0.50 cut-off	≥ 0.85	0.9881	PASS
MQ-4	Brier score	Calibration. The risk tier is derived from the probability, so the probability itself must be truthful	≤ 0.15	0.0458	PASS
	Precision / Recall / Log-loss	Reported alongside for diagnosis	—	0.9518 / 0.9175 / 0.1841	—

Figure 12: Model-quality metrics against their release gates.




Figure 13: ROC, calibration and confusion matrix on the held-out split. The calibration curve tracking the diagonal is what makes the LOW/MEDIUM/HIGH tiering defensible — a 0.70 prediction really does correspond to roughly 70% observed approvals.

Algorithm comparison (carried forward from the Assignment I Analytics Design View)

Candidate	Accuracy	F1	ROC-AUC	Brier	Verdict
Logistic Regression (explainable baseline)	0.9030	0.8789	0.9686	0.0693	Baseline

Candidate	Accuracy	F1	ROC-AUC	Brier	Verdict
Random Forest (chosen)	0.9480	0.9343	0.9881	0.0458	Wins on every metric

Figure 14: The winner is selected automatically by `ModelTrainer.compare_algorithms()` on ROC-AUC — the choice is code, not a comment.

8.2 Data quality — four metrics

In an ML system the data is a dependency exactly as a library is, so it needs its own measurable contract. Two metrics come from the schema gate and two from the drift monitor.

ID	Metric	Definition	Gate	Measured on the 20,000-row training table
DQ-1	Schema conformance rate	Fraction of the 22 declared columns present <i>and</i> within their declared type and range	= 1.00	1.0000 (0 violations)
DQ-2	Missing-value fraction	Overall null ratio, plus the worst single column	≤ 0.02	0.0000 overall; worst column 0.0000
DQ-3	Population Stability Index	Per-feature covariate shift against the frozen training reference profile; the banking-standard drift measure	≤ 0.20	Monitored per batch — see the stress test below
DQ-4	Kolmogorov–Smirnov statistic	Distribution-free two-sample test; catches tail shifts that PSI's coarse bins can miss	$p \geq 0.05$	Reported alongside PSI for every feature

Figure 15: Data-quality metrics. DQ-1 and DQ-2 gate the training run (`strict=True` aborts the build); DQ-3 and DQ-4 run against live batches in production.

Drift stress test

To prove the monitor detects real shift rather than merely running without error, the held-out split was perturbed into a synthetic recession — credit scores down 85 points, debt-to-income $\times 1.6$, card utilisation $\times 1.5$ — and compared against the training reference. The monitor flagged exactly the three perturbed features as SEVERE and left the other 19 STABLE, and emitted a WARNING naming them.

Feature	PSI	KS statistic	KS p-value	Severity
CreditScore	1.7776	0.3669	0.0	SEVERE
DebtToIncomeRatio	0.8469	0.4223	0.0	SEVERE
CreditCardUtilizationRate	0.4618	0.2738	0.0	SEVERE
NetWorth	0.0051	0.0164	0.349087	STABLE
LoanAmount	0.005	0.0243	0.044765	STABLE

Figure 16: Top five features by PSI (reports/metrics/drift_report.csv). The gap between the third and fourth row is the signal: real shift is unambiguous.




Figure 17: PSI per feature against the 0.10/0.25 industry thresholds.




Figure 18: End-to-end scoring latency over 1,000 requests through the real serving path: mean 10.29 ms, p95 17.49 ms, p99 20.04 ms — comfortably inside the 150 ms SLA inherited from Assignment I.

9. Testing in Production, and a Security Consideration

9.1 Approach: shadow deployment, then canary

Offline metrics are necessary and not sufficient. A credit model meets a population it never saw in training, and the label — whether the loan actually defaults — arrives months later. Our release path is therefore staged, and the staging is what makes it safe to be wrong.

Stage	Traffic	What is compared	Promote when	Roll back when
1. Shadow	100% mirrored, 0% served	The candidate scores every live request in parallel with the incumbent; only the incumbent's answer is returned. We diff score distributions, tier mix, disagreement rate and latency.	≥ 2 weeks with no unexplained disagreement and p99 latency within budget	Never — shadow traffic is not served, so a bad candidate cannot harm a customer
2. Canary	5% → 25% → 50% → 100%	Live business KPIs on the canary slice: approval rate, average risk tier, manual-override rate, error rate, p99 latency	Each step held ≥ 48 h with KPIs inside their control limits	Automatic rollback on any breach; the previous artifact is one config change away
3. A/B holdout	Permanent 5% on the incumbent	The metric that actually matters — realised default rate — measured months later against a like-for-like control	Retained permanently as the reference arm	n/a

Figure 19: Staged rollout. Shadow answers 'does it behave?'; canary answers 'does it behave on customers?'; the A/B holdout answers 'was it actually better?'.

Why shadow first for this system specifically. The outcome label is delayed by months, so we cannot A/B-test our way to a decision quickly. Shadow mode gives an immediate, zero-risk read on the one thing that is observable on day one — whether the candidate's score distribution and its disagreements with the incumbent are explainable. The infrastructure for this already exists in the codebase: `ModelRegistry` makes 'which artifact is loaded' a configuration value, and every scoring event is already logged as JSON with `model_version`, so comparing two arms is a log query rather than new instrumentation. `DriftMonitor` supplies the population check for the same window.

9.2 Security consideration: the input boundary as an attack surface

The threat we treat as primary is **adversarial and malformed input at the scoring endpoint**. The endpoint is reachable from a customer-facing portal, it accepts a 20-field numeric payload, and its output is a lending decision with direct financial consequence. That combination invites three concrete abuses: gaming the model by probing which field flips a decision; feeding out-of-distribution values to push the model into a region where its behaviour is untested; and classic injection or resource exhaustion against the service itself.

Control	Where	Effect
Bounded field constraints on all 20 inputs	api/schemas.py	Values outside the legitimate domain (credit score 1500, income 10 ¹²) are rejected with 422 before any model code runs — the estimator is never asked to extrapolate
extra="forbid"	api/schemas.py	Unknown keys are rejected instead of ignored; blocks parameter-pollution probing and catches client typos
Strict type coercion	Pydantic v2	A SQL/JS injection string in a numeric field fails int_parsing — the payload never becomes a query or a template
Business-rule filter	RiskPredictor.validate_business_rules	A second, semantic gate behind the syntactic one: schema-valid but economically absurd applications are refused with 400
Batch size cap (100)	BatchLoanApplications	Bounds the work one request can demand — a basic DoS control
Opaque error envelope	Exception handlers in api/app.py	Clients get a reason, never a stack trace, module path or internal identifier
Full audit log	logging_utils.py	Every decision is recorded with its inputs, probability, tier and model_version, so probing patterns are detectable after the fact and every decision is reconstructable for a regulator

Figure 20: Layered input-validation controls. Six integration tests parametrised over injection strings, oversized batches and extreme numerics assert that each control holds.

What this does not solve. Validation constrains each field independently; it cannot detect an application in which every value is individually plausible but the combination is fabricated. The honest mitigations for that are outside the request path: corroborating declared income and liabilities against the bureau feed rather than trusting the payload, rate-limiting and authenticating per client so systematic probing is visible and attributable, and monitoring the score distribution per caller for the signature of an optimisation attack. Model access control — the artifact and the reference profile are deployment assets, not public ones — belongs in the same list.

Appendix A — Repository Layout

```

Group_84/
■■■ configs/config.yaml           # thresholds, gates, hyper-parameters, feature contract
■■■ data/
■   ■■■ generate_synthetic_data.py # schema-faithful stand-in for the Kaggle file
■   ■■■ prepare_data.py           # cleaning, encoding, feature engineering
■   ■■■ loan_data_processed.csv   # 20,000 x 23 modelling table
■■■ src/loan_risk/
■   ■■■ config.py                 # frozen dataclasses loaded from YAML
■   ■■■ exceptions.py             # domain exception hierarchy
■   ■■■ logging_utils.py          # structured JSON logging + Timer
■   ■■■ data/ingestion.py         # DataIngestor
■   ■■■ data/validation.py        # DataValidator -> DQ-1, DQ-2
■   ■■■ features/engineering.py   # FeatureEngineer (sklearn transformer)
■   ■■■ models/trainer.py         # ModelTrainer + quality gates
■   ■■■ models/predictor.py       # ModelRegistry + RiskPredictor
■   ■■■ monitoring/drift.py       # DriftMonitor -> DQ-3, DQ-4
■   ■■■ api/{schemas,app}.py      # FastAPI contract and routes
■■■ tests/                        # 84 tests: unit | integration | data | ml
■■■ scripts/                      # train, evaluate, render evidence, build report
■■■ notebooks/research_prototype.ipynb # RESEARCH CODE (evidence)
■■■ legacy/research_feature_prototype.py # same, exported for the linters
■■■ reports/{lint,metrics,figures}/  # all captured evidence
■■■ artifacts/model.joblib         # serialised pipeline

```

Appendix B — Reproducing Every Number in This Report

```

pip install -r requirements.txt

python data/generate_synthetic_data.py --rows 20000 # raw table
python data/prepare_data.py                       # 22 features + target
python scripts/train_model.py                     # train, gate, persist
python -m pytest tests -v                         # 84 tests
python scripts/evaluate_and_report.py              # metrics + figures

python -m isort src scripts tests
python -m black src scripts tests
python -m flake8 src scripts tests
python -m pylint src/loan_risk

PYTHONPATH=src python -m uvicorn loan_risk.api.app:app --port 8000
# -> http://127.0.0.1:8000/docs (generated Swagger UI)

```

Appendix C — Full Test Inventory

All 84 tests, as collected by pytest.

Unit tests — test_unit_features.py (14)

```

test_safe_ratio_matches_expected_value
test_safe_ratio_never_divides_by_zero
test_compute_loan_to_income_is_monotonic_in_loan_amount
test_compute_savings_to_loan_is_monotonic_in_savings
test_transformer_emits_the_declared_feature_contract
test_transformer_is_deterministic
test_transformer_does_not_mutate_its_input
test_transformer_recomputes_derived_columns_from_base_columns
test_transformer_rejects_missing_base_columns
test_transformer_rejects_non_dataframe_input
test_transformer_raises_on_non_finite_derived_values
test_settings_expose_a_non_empty_feature_contract
test_settings_are_immutable
test_missing_config_file_fails_loudly

```

Integration tests — test_integration_api.py (21)

```

test_health_endpoint_reports_a_loaded_model
test_model_metadata_endpoint_exposes_the_contract
test_openapi_schema_is_generated
test_predict_returns_200_and_a_complete_payload
test_batch_predict_returns_one_result_per_application
test_schema_violation_returns_422
test_missing_required_field_returns_422
test_unknown_field_returns_422
test_wrong_type_returns_422
test_business_rule_violation_returns_400
test_oversized_batch_returns_422
test_empty_batch_returns_422
test_unknown_route_returns_404
test_adversarial_payloads_are_rejected_at_the_boundary['; DROP TABLE applications;--]
test_adversarial_payloads_are_rejected_at_the_boundary[<script>alert(1)</script>]
test_adversarial_payloads_are_rejected_at_the_boundary[../../../../etc/passwd]
test_adversarial_payloads_are_rejected_at_the_boundary[malicious3]
test_adversarial_payloads_are_rejected_at_the_boundary[malicious4]
test_extreme_numeric_values_are_rejected
test_error_responses_never_leak_internals
test_api_and_direct_pipeline_agree

```

Data-validation tests — test_data_validation.py (15)

```

test_clean_data_conforms_fully
test_out_of_range_credit_score_is_flagged
test_negative_income_is_flagged
test_missing_column_is_flagged
test_strict_mode_raises_on_violation
test_validator_rejects_empty_frame
test_every_declared_column_has_a_spec
test_missing_value_fraction_is_measured
test_psi_is_near_zero_for_identical_distributions
test_psi_detects_a_shifted_distribution
test_psi_rejects_empty_samples
test_drift_monitor_flags_only_the_shifted_feature
test_ingestor_rejects_a_missing_file
test_ingestor_rejects_a_frame_missing_contract_columns
test_ingestor_split_returns_contract_ordered_features

```

ML tests — training — test_model_training.py (11)

```

test_model_can_overfit_a_small_batch
test_training_loss_decreases_with_capacity
test_training_is_reproducible
test_shuffled_labels_destroy_generalisation
test_real_labels_beat_chance
test_quality_gates_reject_a_weak_model
test_training_rejects_a_single_class_target
test_training_rejects_length_mismatch
test_unknown_algorithm_is_rejected
test_model_quality_metrics_meet_release_gates
test_artifact_round_trips_through_disk

```

ML tests — inference — test_model_inference.py (23)

```

test_batch_prediction_shape_is_correct
test_probabilities_are_valid_and_sum_to_one
test_single_prediction_returns_a_well_formed_assessment
test_output_stays_in_range_across_the_credit_spectrum[300]
test_output_stays_in_range_across_the_credit_spectrum[500]
test_output_stays_in_range_across_the_credit_spectrum[700]
test_output_stays_in_range_across_the_credit_spectrum[850]
test_higher_credit_score_does_not_reduce_approval_probability
test_higher_debt_to_income_does_not_increase_approval_probability
test_a_much_larger_loan_does_not_increase_approval_probability
test_prediction_is_invariant_to_dictionary_key_order
test_prediction_is_invariant_to_repeated_calls
test_prediction_is_invariant_to_batching
test_prediction_is_invariant_to_column_order
test_previous_default_forces_the_high_risk_tier
test_risk_tier_boundaries[0.95-LOW]
test_risk_tier_boundaries[0.55-MEDIUM]
test_risk_tier_boundaries[0.1-HIGH]
test_business_rule_rejects_over_leveraged_application
test_business_rule_rejects_sub_threshold_income
test_validation_filter_does_not_mutate_the_payload
test_inference_without_a_loaded_model_raises
test_inference_latency_is_within_sla

```